
High-Performance Implementation of the Fast Fourier Transform Using the Cooley–Tukey Algorithm

Sharvari Satish Deshmukh
A59022896

ECE 277 - Fall 2024

Discrete Fourier Transform & Fast Fourier Transform

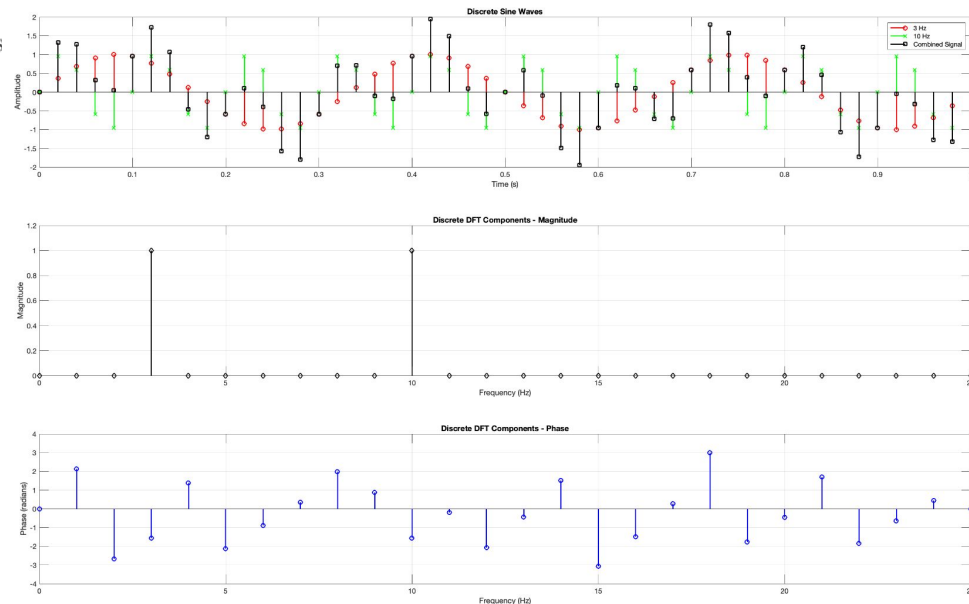
Let x_0, \dots, x_{n-1} be **complex numbers**. The **DFT** is defined as

$$X_k = \sum_{m=0}^{n-1} x_m e^{-i \frac{2\pi km}{n}}, \quad k = 0, \dots, n-1,$$

where $e^{i \frac{2\pi}{n}}$ is a **primitive** n th root of 1.

DFT ----> FFT

$O(n^2)$ -----> $O(n \log_2(n))$



Cooley-Tukey: Recursive vs Iterative Algorithm

```
def FFT_recursive(P):  
    # P = [p0, p1, ..., pn-1] coeff representation  
    n = len(P) # n is a power of 2  
    if n == 1:  
        return P  
     $\omega = e^{(2\pi i/n)}$   
    Pe, Po = [p0, p2, ..., pn-2], [p1, p3, ..., pn-1]  
    ye, yo = FFT(Pe), FFT(Po)  
    y = [0] * n  
    for j in range(n/2):  
        y[j] = ye[j] +  $\omega^j$  * yo[j]  
        y[j + n/2] = ye[j] -  $\omega^j$  * yo[j]  
    return y
```

```
def FFT_iterative(P):  
    # P = [p0, p1, ..., pn-1] coeff representation  
    n = len(P) # n is a power of 2  
    log_n = log2(n)  
    # Bit-reversal permutation  
    P = bit_reverse_copy(P)  
    for s in range(1, log_n + 1):  
        m = 2^s  
         $\omega_m = e^{(2\pi i/m)}$   
        for k in range(0, n, m):  
             $\omega = 1$   
            for j in range(m//2):  
                t =  $\omega$  * P[k + j + m//2]  
                u = P[k + j]  
                P[k + j] = u + t  
                P[k + j + m//2] = u - t  
                 $\omega = \omega * \omega_m$   
    return P
```

Coolley-Tukey: Iterative (Radix-2)

```
def FFT_iterative(P):
    # P - [p0, p1, ..., pn-1] coeff representation
    n = len(P) # n is a power of 2
    log_n = log2(n)
    # Bit-reversal permutation
    P = bit_reverse_copy(P)
    for s in range(1, log_n + 1):
        m = 2^s
        w_m = e^(2pi/m) # twiddle factor
        for k in range(0, n, m):
            w = 1
            for j in range(m//2):
                t = w * P[k + j + m//2]
                u = P[k + j]
                P[k + j] = u + t
                P[k + j + m//2] = u - t
                w = w * w_m
            # i=3 (011) --> (110) i=6
            # => swap(3,6)
        # swap(3,6)
    return P
```






Figure 1: Length-8 Radix-2 FFT Flow Graph

Simple GPU implementation

Nvidia RTX A2000 - Execution Time



```
__global__ void fft_stage_kernel(float2* d_data, int N, int s)
{
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    int halfSize = 1 << (s - 1); // m2
    int fftSize = 1 << s; // m

    if (tid < N / 2) {
        int group = tid / halfSize;
        int j = tid % halfSize;
        int k = group * fftSize;

        float angle = -2.0f * (float)M_PI * j / (float)fftSize;
        float2 w = make_float2(cosf(angle), sinf(angle));

        int index1 = k + j;
        int index2 = k + j + halfSize;


        float2 u = d_data[index1];
        float2 t = d_data[index2];

        // Complex multiplication t * w
        float2 temp;
        temp.x = w.x * t.x - w.y * t.y;
        temp.y = w.x * t.y + w.y * t.x;

        // Butterfly
        d_data[index1].x = u.x + temp.x;
        d_data[index1].y = u.y + temp.y;
        d_data[index2].x = u.x - temp.x;
        d_data[index2].y = u.y - temp.y;
    }
}
```

Optimized GPU: Streams & Shared Memory

Nvidia RTX A2000 - Execution Time



```
__global__ void fft_kernel_sm_all_stages(float2* d_data, int N, int logN)
{
    extern __shared__ float2 s_data[];
    int tid = threadIdx.x;
    s_data[tid] = d_data[tid];

    for (int s = 1; s <= logN; s++) {
        int m = 1 << s;
        int m2 = m >> 1;

        __syncthreads();

        if (tid < N / 2) {
            int group = tid / m2;
            int j = tid % m2;
            int k = group * m;

            float angle = -2.0f * (float)M_PI * (float)j / (float)m;
            float2 w = make_float2(cosf(angle), sinf(angle));

            int i1 = k + j;
            int i2 = k + j + m2;

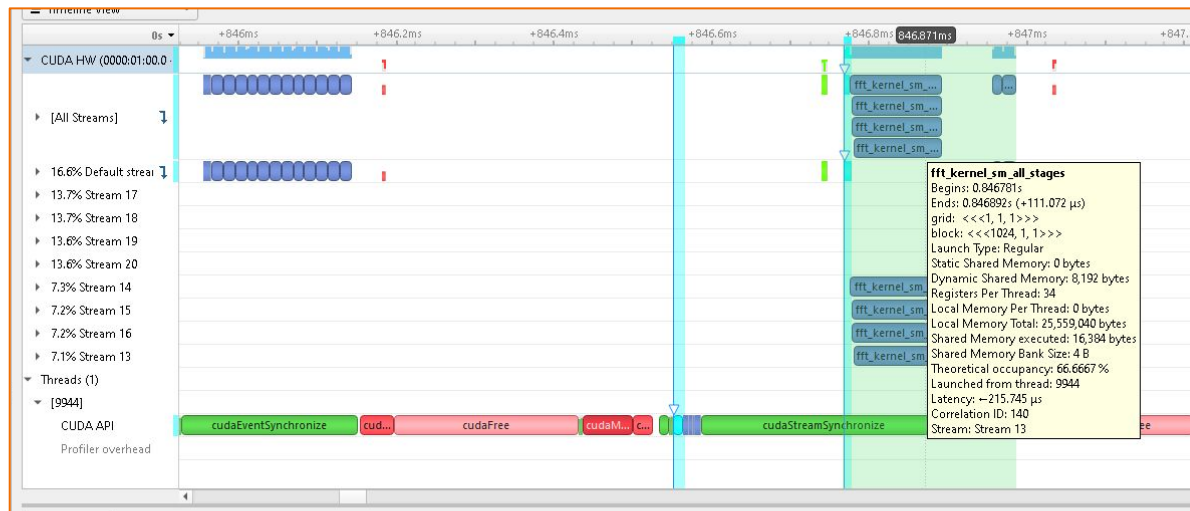
            float2 u = s_data[i1];
            float2 t = s_data[i2];

            // Complex multiplication t * w
            float2 temp;
            temp.x = w.x * t.x - w.y * t.y;
            temp.y = w.x * t.y + w.y * t.x;

            // Butterfly
            s_data[i1].x = u.x + temp.x;
            s_data[i1].y = u.y + temp.y;
            s_data[i2].x = u.x - temp.x;
            s_data[i2].y = u.y - temp.y;
        }
    }

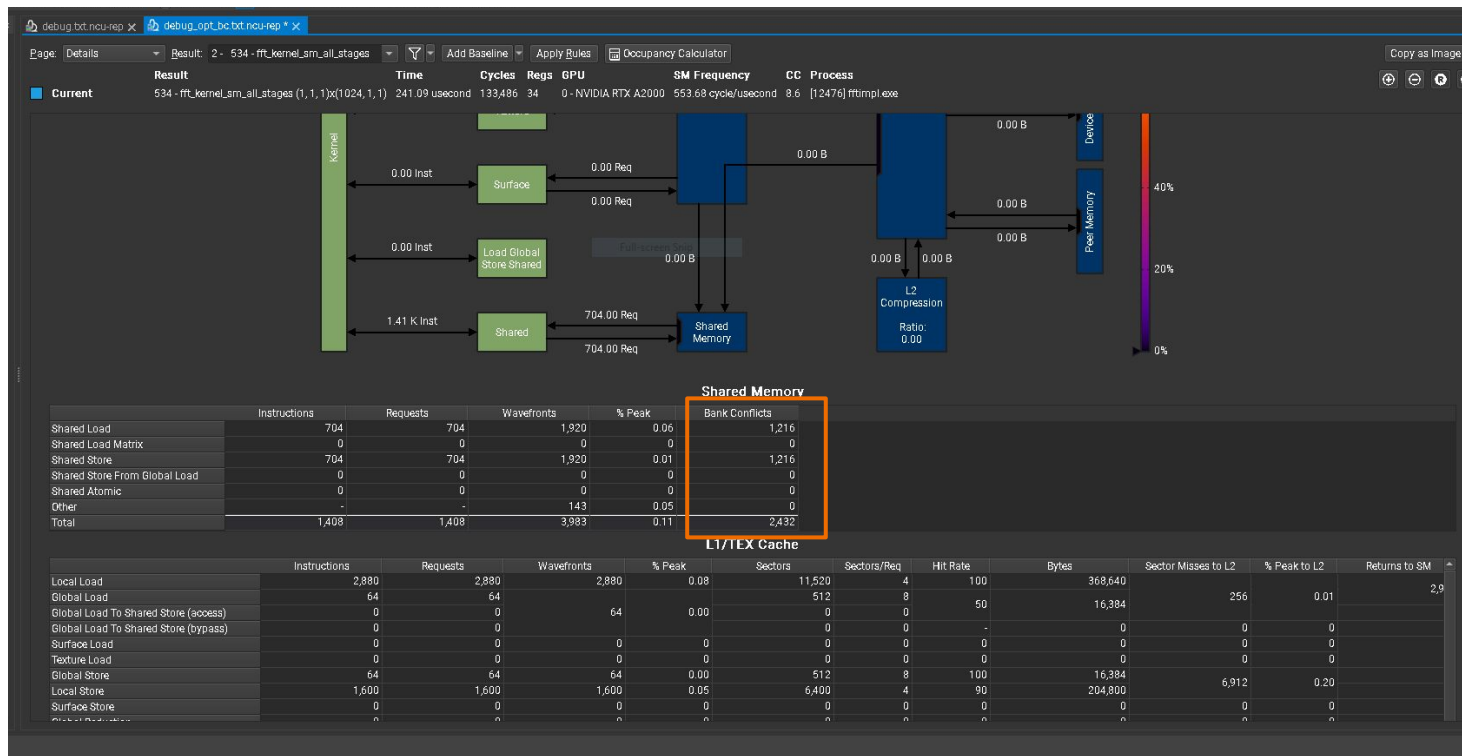
    __syncthreads();
    d_data[tid] = s_data[tid];
}
```

Optimized GPU: Streams & Shared Memory




N = 4096

Optimized GPU: Bank Conflicts



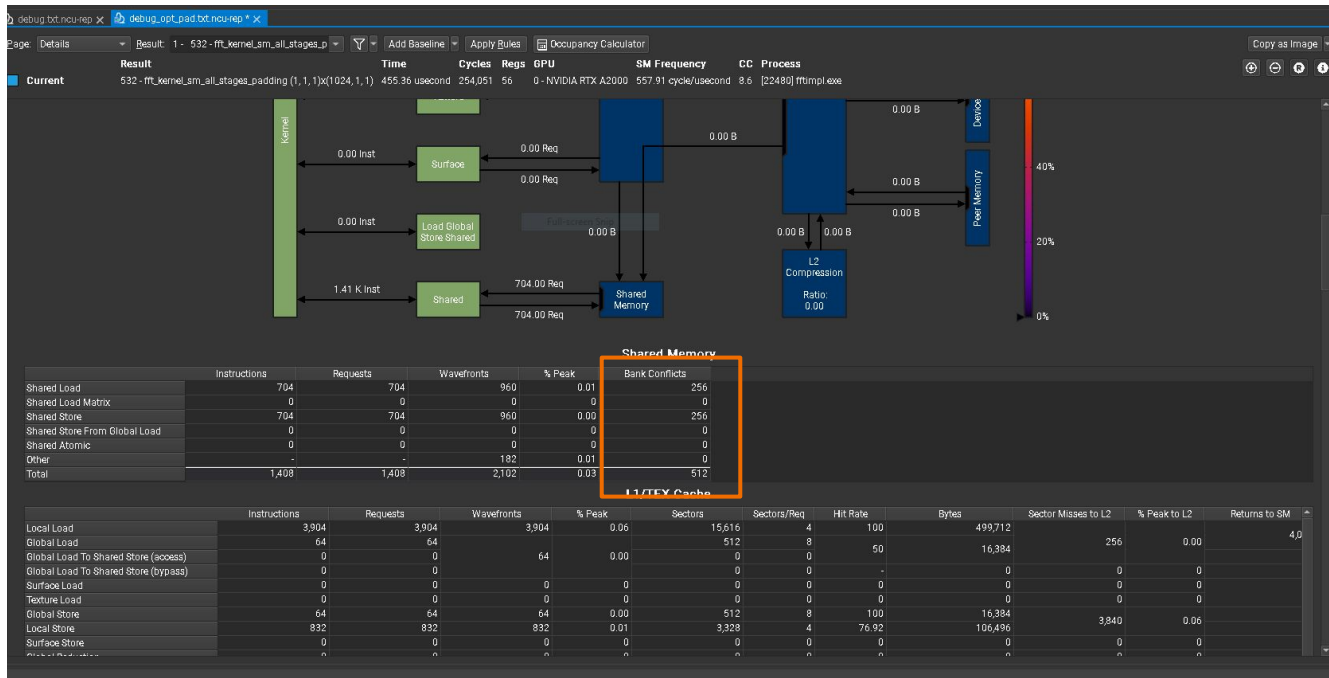
Optimized GPU: Padding

Nvidia RTX A2000 - Execution Time



```
__global__ void fft_kernel_sm_all_stages(float2* d_data, int N, int logN)
{
    extern __shared__ float s_mem[];
    const int warpSize = 32;
    const int paddingPerWarp = 4;
    const int numWarps = N / warpSize;
    const int paddedLength = N + numWarps * paddingPerWarp;
    float* s_real = s_mem;
    float* s_imag = s_mem + paddedLength;
    int tid = threadIdx.x;
    auto paddedIndex = [=](int i) {
        int warpId = i / warpSize;
        int offset = warpId * paddingPerWarp;
        return i + offset;
    };
    float2 val = d_data[tid];
    s_real[paddedIndex(tid)] = val.x;
    s_imag[paddedIndex(tid)] = val.y;
    for (int s = 1; s <= logN; s++) {
        int m = 1 << s;
        int m2 = m >> 1;
        __syncthreads();
        if (tid < N / 2) {
            int group = tid / m2;
            int j = tid % m2;
            int k = group * m;
            float angle = -2.0f * (float)M_PI * (float)j / (float)m;
            float2 w = make_float2(cosf(angle), sinf(angle));
            int i1 = k + j;
            int i2 = k + j + m2;
            float ur = s_real[paddedIndex(i1)];
            float ui = s_imag[paddedIndex(i1)];
            float tr = s_real[paddedIndex(i2)];
            float ti = s_imag[paddedIndex(i2)];
            float temp_r = w.x * tr - w.y * ti;
            float temp_i = w.x * ti + w.y * tr;
            s_real[paddedIndex(i1)] = ur + temp_r;
            s_imag[paddedIndex(i1)] = ui + temp_i;
            s_real[paddedIndex(i2)] = ur - temp_r;
            s_imag[paddedIndex(i2)] = ui - temp_i;
        }
        __syncthreads();
        val.x = s_real[paddedIndex(tid)];
        val.y = s_imag[paddedIndex(tid)];
        d_data[tid] = val;
    }
}
```

Optimized GPU: Padding



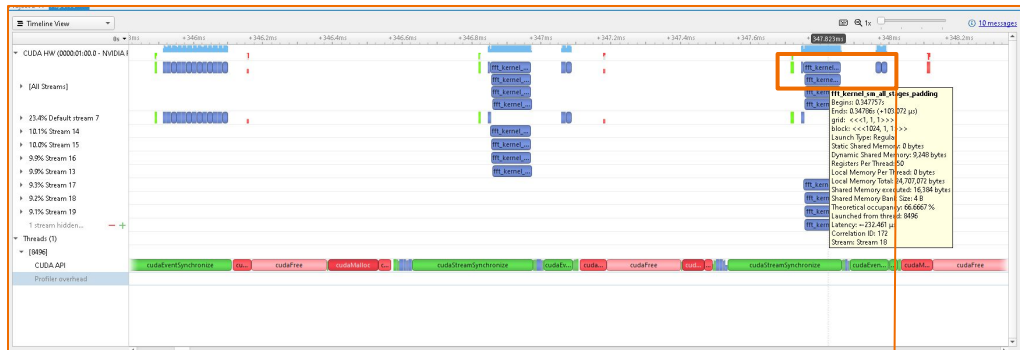
Optimized GPU: Instructions Overhead

```
__global__ void fft_kernel_sm_all_stages(float2* d_data, int N, int logN)
{
    extern __shared__ float s_mem[];
    const int warpSize = 32;
    const int paddingPerWarp = 4;
    const int numWarps = N / warpSize;
    const int paddedLength = N + numWarps * paddingPerWarp;
    float* s_real = s_mem;
    float* s_imag = s_mem + paddedLength;
    int tid = threadIdx.x;
    auto paddedIndex = [=](int i) {
        int warpId = i / warpSize;
        int offset = warpId * paddingPerWarp;
        return i + offset;
    };
    float2 val = d_data[tid];
    s_real[paddedIndex(tid)] = val.x;
    s_imag[paddedIndex(tid)] = val.y;
    for (int s = 1; s <= logN; s++) {
        int m = 1 << s;
        int m2 = m >> 1;
        __syncthreads();
        if (tid < N / 2) {
            int group = tid / m2;
            int j = tid % m2;
            int k = group * m;
            float angle = -2.0f * (float)M_PI * (float)j / (float)m;
            float2 w = make_float2(cosf(angle), sinf(angle));
            int i1 = k + j;
            int i2 = k + j + m2;
            float ur = s_real[paddedIndex(i1)];
            float ui = s_imag[paddedIndex(i1)];
            float tr = s_real[paddedIndex(i2)];
            float ti = s_imag[paddedIndex(i2)];
            float temp_r = w.x * tr - w.y * ti;
            float temp_i = w.x * ti + w.y * tr;
            s_real[paddedIndex(i1)] = ur + temp_r;
            s_imag[paddedIndex(i1)] = ui + temp_i;
            s_real[paddedIndex(i2)] = ur - temp_r;
            s_imag[paddedIndex(i2)] = ui - temp_i;
        }
        __syncthreads();
        val.x = s_real[paddedIndex(tid)];
        val.y = s_imag[paddedIndex(tid)];
        d_data[tid] = val;
    }
}
```

```
__global__ void fft_kernel_sm_all_stages_padding(float2* d_data, int N, int logN)
{
    extern __shared__ float s_mem[];

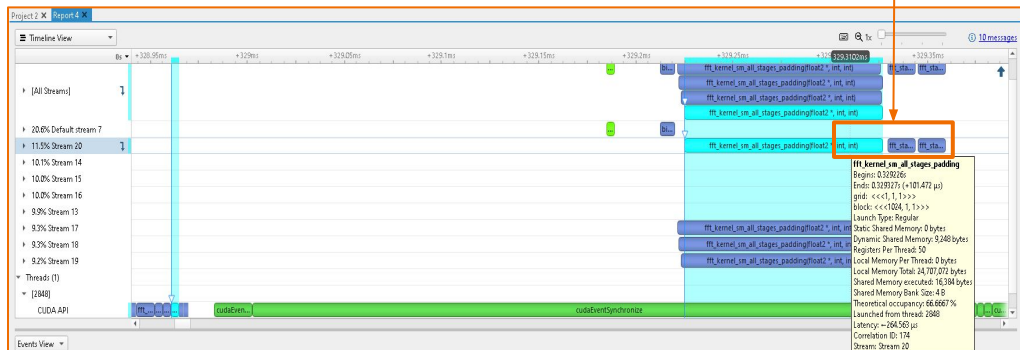
    const int tid = threadIdx.x;
    const int warpSize = WARP_SIZE;
    const int paddingPerWarp = PADDING_SIZE;
    const int numWarps = N / warpSize;
    const int paddedLength = N + numWarps * paddingPerWarp;
    float* s_real = s_mem;
    float* s_imag = s_mem + paddedLength;
    int warpId_t = tid / warpSize;
    int pTid = tid + warpId_t * paddingPerWarp;
    float2 val = d_data[tid];
    s_real[pTid] = val.x;
    s_imag[pTid] = val.y;
    for (int s = 1; s <= logN; s++) {
        int m = 1 << s;
        int m2 = m >> 1;
        float angleBase = -2.0f * (float)M_PI / (float)m;
        __syncthreads();
        if (tid < N / 2) {
            int group = tid / m2;
            int j = tid % m2;
            int k = group * m;
            float angle = angleBase * j;
            float2 w = make_float2(__cosf(angle), __sinf(angle));
            int i1 = k + j;
            int i2 = k + j + m2;
            int warpId_i1 = i1 / warpSize;
            int pI1 = i1 + warpId_i1 * paddingPerWarp;
            int warpId_i2 = i2 / warpSize;
            int pI2 = i2 + warpId_i2 * paddingPerWarp;
            float ur = s_real[pI1];
            float ui = s_imag[pI1];
            float tr = s_real[pI2];
            float ti = s_imag[pI2];
            float temp_r = w.x * tr - w.y * ti;
            float temp_i = w.x * ti + w.y * tr;
            s_real[pI1] = ur + temp_r;
            s_imag[pI1] = ui + temp_i;
            s_real[pI2] = ur - temp_r;
            s_imag[pI2] = ui - temp_i;
        }
        __syncthreads();
        val.x = s_real[pTid];
        val.y = s_imag[pTid];
        d_data[tid] = val;
    }
}
```

Optimized GPU: Synchronization Overhead



```
for (int i = 0; i < numSegments; i++)
    cudaStreamSynchronize(streams[i]);

if (N > baseSize) {
    for (int s = logBaseSize + 1; s <= logN; s++) {
        int m = 1 << s;
        int totalPairs = N / 2;
        int gridSize = (totalPairs + blockSize - 1) / blockSize;
        fft_stage_kernel << <gridSize, blockSize, 0>> (d_data, N, s);
    }
}
```



```
// for (int i = 0; i < numSegments; i++)
//     cudaStreamSynchronize(streams[i]);

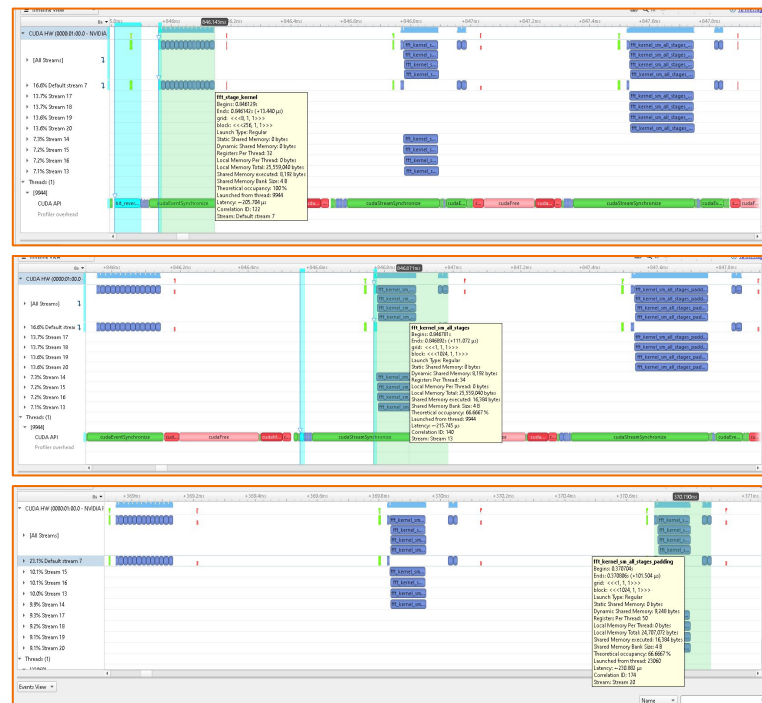

if (N > baseSize) {
    for (int s = logBaseSize + 1; s <= logN; s++) {
        int m = 1 << s;
        int totalPairs = N / 2;
        int gridSize = (totalPairs + blockSize - 1) / blockSize;
        fft_stage_kernel << <gridSize, blockSize, 0>> (d_data, N, s);
    }
}
```

N = 4096

DEMO

Results

Nvidia RTX A2000 - Execution Time



N = 4096

Limitations & Future Scope


- Algorithmic limitations:
 - Dependence on Input size
 - Sequential dependency
 - Butterfly design
- Shared memory size limitation
- Synchronization overhead
- Floating point data:
 - Precision vs Performance
 - Cascade of error
- Hybrid algorithms, zero-padding (Radix-4/8, Stockham Algorithm)
- Mixed precision (Tensor Cores), error-compensating algorithms
- Batch Processing, Hardware-Aware Optimizations
- State-of-the-art FFT Libraries

References

- [1] Burrus, Sidney. "The Cooley-Tukey Fast Fourier Transform Algorithm * C ." (2014).
- [2] https://www.wikiwand.com/en/articles/Fast_Fourier_transform
- [3] https://www.youtube.com/watch?v=h7ap07q16V0&t=1336s&ab_channel=Reducible
- [4] <https://github.com/KAdamek/SMFFT>
- [5] <https://forums.developer.nvidia.com/t/does-cufft-show-much-higher-efficiency-than-cpu-fft-routines/17790/4>
- [6] <https://github.com/roguh/cuda-fft/tree/main>
- [7] <https://github.com/anair-eng/CUDA-MPI-pthreads-FFT>
- [8] <https://cs.wmich.edu/gupta/teaching/cs5260/5260Sp15web/studentProjects/tiba&hussein/03278999.pdf>

Thank You!

Appendix

Nvidia RTX A2000						
N	blockSize	CPU	Naive GPU	Optimized GPU	Ultra Optimized	
32	256	0.0089	0.0768	0.155648	0.16384	
64	256	0.0118	0.090112	0.149504	0.14336	
128	256	0.0178	0.1024	0.151552	0.141312	
256	256	0.0301	0.114688	0.149536	0.144384	
512	256	0.06	0.129024	0.161888	0.142336	
1024	256	0.1254	0.154624	0.150528	0.14336	
2048	256	0.2513	0.156672	0.164864	0.156672	
4096	256	0.5649	0.172032	0.195584	0.171008	
8192	256	1.2737	0.186368	0.202752	0.191488	
16384	256	2.6371	0.205824	0.246784	0.2456	
32768	256	5.7456	0.235552	0.394368	0.395264	
65536	256	12.3248	0.449536	0.63376	0.579584	